

# TensorInference: A Julia package for tensor-based probabilistic inference

Martin Roa-Villescas<sup>1\*</sup> and Jin-Guo Liu<sup>2\*</sup>

<sup>1</sup> Eindhoven University of Technology <sup>2</sup> Hong Kong University of Science and Technology (Guangzhou)

\* These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

## Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright  
and release the work under a  
Creative Commons Attribution 4.0  
International License ([CC BY 4.0](#)).

## Summary

TensorInference.jl is a Julia ([Bezanson et al., 2017](#)) package designed for performing probabilistic inference in discrete graphical models. Capitalizing on the recent advances in the field of tensor networks ([Orús, 2014, 2019](#); [Robeva & Seigal, 2019](#)), TensorInference.jl offers high-performance solutions for prevalent inference problems. Specifically, it provides methods to:

1. calculate the partition function (also known as the probability of evidence).
2. compute the marginal probability distribution over each variable given evidence.
3. find the most likely assignment to all variables given evidence.
4. find the most likely assignment to a set of query variables after marginalizing out the remaining variables.
5. draw samples from the posterior distribution given evidence ([Cheng et al., 2019](#); [Han et al., 2018](#)).

The use of a tensor network-based infrastructure offers several advantages when dealing with complex computational tasks. Firstly, it simplifies the process of computing gradients by employing differentiable programming ([Liao et al., 2019](#)), a critical operation for the aforementioned inference tasks. Secondly, it supports generic element types without a significant compromise on performance. The advantage of supporting generic element types lies in the ability to solve a variety of problems using the same tensor network contraction algorithm, simply by varying the element types used. This flexibility has allowed us to seamlessly implement solutions for several of the inference tasks mentioned earlier ([Jin Guo Liu et al., 2022](#); [Jin-Guo Liu et al., 2021](#)). Thirdly, it allows users to define a hyper-optimized contraction order, which is known to have a significant impact on the computational performance of contracting tensor networks ([Gao et al., 2021](#); [Markov & Shi, 2008](#); [Pan & Zhang, 2021](#)). TensorInference.jl provides a predefined set of state-of-the-art contraction ordering methods. These methods include a *local search based method* (TreeSA) ([Kalachev et al., 2022](#)), two *min-cut based methods* (KaHyParBipartite) ([Gray & Kourtis, 2021](#)) and (SABipartite), and a *greedy method* (GreedyMethod). Finally, tensor networks – and by extension, TensorInference.jl – harness the latest developments in computational technology, including a highly optimized set of BLAS routines ([Blackford et al., 2002](#)) and GPU technology.

## Statement of need

A major challenge in developing intelligent systems is the ability to reason under uncertainty, a challenge that appears in many real-world problems across various domains, including artificial intelligence, medical diagnosis, computer vision, computational biology, and natural language processing. Reasoning under uncertainty involves calculating the probabilities of relevant variables while taking into account any information that is acquired. This process, which can

be thought of as drawing global insights from local observations, is known as *probabilistic inference*.

*Probabilistic graphical models* (PGMs) provide a unified framework to perform probabilistic inference. These models use graphs to represent the joint probability distribution of complex systems in a concise manner by exploiting the conditional independence between variables in the model. Additionally, they form the foundation for various algorithms that enable efficient probabilistic inference.


However, even with the representational aid of PGMs, performing probabilistic inference remains an intractable endeavor on many real-world models. The reason is that performing probabilistic inference involves complex combinatorial optimization problems in very high dimensional spaces. To tackle these challenges, more efficient and scalable inference algorithms are needed.

As an attempt to tackle the aforementioned challenges, we present `TensorInference.jl`, a Julia package for probabilistic inference that combines the representational capabilities of PGMs with the computational power of tensor networks. By harnessing the best of both worlds, `TensorInference.jl` aims to enhance the performance of probabilistic inference, thereby expanding the tractability spectrum of exact inference for more complex, real-world models.

`TensorInference.jl` succeeds `JunctionTrees.jl` (Roa-Villescas et al., 2022, 2023), a Julia package implementing the Junction Tree Algorithm (JTA) (Jensen et al., 1990; Lauritzen & Spiegelhalter, 1988). While the latter employs tensor-based technology to optimize the computation of individual sum-product messages within the JTA context, `TensorInference.jl` takes a different route. It adopts a holistic tensor network approach, which opens new doors for optimization opportunities, and significantly reduces the algorithm's complexity compared to the JTA.

## Usage example

The graph below corresponds to the *ASIA network* (Lauritzen & Spiegelhalter, 1988), a simple Bayesian network (Pearl, 1985) used extensively in educational settings. It describes the probabilistic relationships between different random variables which correspond to possible diseases, symptoms, risk factors and test results.



Random variable	Meaning
A	Recent trip to Asia
T	Patient has tuberculosis
S	Patient is a smoker
L	Patient has lung cancer
B	Patient has bronchitis
E	Patient has T and/or L
X	Chest X-Ray is positive
D	Patient has dyspnoea

**Figure 1:** The ASIA network: a simplified example of a Bayesian network from the context of medical diagnosis (Lauritzen & Spiegelhalter, 1988).

In the example, a patient has recently visited Asia and is now experiencing dyspnea. These conditions serve as the evidence for the observed variables ( $A$  and  $D$ ). The doctor's task is to assess the likelihood of various diseases — tuberculosis, lung cancer, and bronchitis — which constitute the query variables in this scenario ( $T$ ,  $L$ , and  $B$ ).

We now demonstrate how to use `TensorInference.jl` for conducting a variety of inference

75 tasks on this toy example. Please note that as the API may evolve, we recommend checking  
76 the [examples](#) directory of the official `TensorInference.jl` repository for the most up-to-date  
77 version of this example.

```
# Import the TensorInference package, which provides the functionality needed  
# for working with tensor networks and probabilistic graphical models.  
using TensorInference  
  
# Load the ASIA network model from the `asia.uai` file located in the examples  
# directory. Refer to the documentation of this package for a description of the  
# format of this file.  
instance = read_instance(pkgdir(TensorInference), "examples", "asia", "asia.uai")  
  
# Create a tensor network representation of the loaded model.  
# The variable 7 is the variable of interest, which will be retained in the output.  
tn = TensorNetworkModel(instance; openvars=[7])  
  
# Calculate the partition function for each assignment of variable 7.  
probability(tn)  
  
# Calculate the marginal probabilities of each random variable in the model.  
marginals(tn)  
  
# Retrieve the variables associated with the tensor network model.  
get_vars(tn)  
  
# Assume that the "X-ray" result (variable 7) is positive.  
# Since setting an evidence may affect the contraction order of the tensor  
# network, recompute it.  
tn = TensorNetworkModel(instance; evidence=Dict{7 => 0})  
  
# Calculate the maximum log-probability among all configurations.  
maximum_logp(tn)  
  
# Generate 10 samples from the probability distribution represented by the  
# model.  
sample(tn, 10)  
  
# Retrieve both the maximum log-probability and the most probable  
# configuration. In this configuration, the most likely outcomes are that the  
# patient smokes (variable 3) and has lung cancer (variable 4).  
logp, cfg = most_probable_config(tn)  
  
# Compute the most probable values of certain variables (e.g., 4 and 7) while  
# marginalizing over others. This is known as Maximum a Posteriori (MAP)  
# estimation.  
mmap = MMAPModel(instance, queryvars=[4, 7])  
  
# Get the most probable configurations for variables 4 and 7.  
most_probable_config(mmap)  
  
# Compute the total log-probability of having lung cancer. The results suggest  
# that the probability is roughly half.  
log_probability(mmap, [1, 0]), log_probability(mmap, [0, 0])
```

## Acknowledgments

This work is partially funded by the Netherlands Organization for Scientific Research and the Guangzhou Municipal Science and Technology Project (No. 2023A03J0003). We extend our gratitude to Madelyn Cain and Patrick Wijnings for their insightful discussions on the intersection of tensor networks and probabilistic graphical models.

## References

- Bezanson, J., Edelman, A., Karpinski, S., & Shah, V. (2017). Julia: A Fresh Approach to Numerical Computing. *SIAM Review*, 59(1), 65–98. <https://doi.org/10.1137/141000671>
- Blackford, L. S., Petitet, A., Pozo, R., Remington, K., Whaley, R. C., Demmel, J., Dongarra, J., Duff, I., Hammarling, S., Henry, G., & others. (2002). An updated set of basic linear algebra subprograms (BLAS). *ACM Transactions on Mathematical Software*, 28(2), 135–151.
- Cheng, S., Wang, L., Xiang, T., & Zhang, P. (2019). Tree tensor networks for generative modeling. *Physical Review B*, 99(15), 155131.
- Gao, X., Kalinowski, M., Chou, C.-N., Lukin, M. D., Barak, B., & Choi, S. (2021). Limitations of linear cross-entropy as a measure for quantum advantage. *arXiv Preprint arXiv:2112.01657*.
- Gray, J., & Kourtis, S. (2021). Hyper-optimized tensor network contraction. *Quantum*, 5, 410. <https://doi.org/10.22331/q-2021-03-15-410>
- Han, Z.-Y., Wang, J., Fan, H., Wang, L., & Zhang, P. (2018). Unsupervised generative modeling using matrix product states. *Physical Review X*, 8(3), 031012.
- Jensen, F. V., Lauritzen, S. L., & Olesen, K. G. (1990). Bayesian updating in causal probabilistic networks by local computations. *Computational Statistics Quarterly*, 4, 269–282.
- Kalachev, G., Panteleev, P., & Yung, M.-H. (2022). *Multi-tensor contraction for XEB verification of quantum circuits*. <https://arxiv.org/abs/2108.05665>
- Lauritzen, S. L., & Spiegelhalter, D. J. (1988). Local computations with probabilities on graphical structures and their application to expert systems. *Journal of the Royal Statistical Society: Series B (Methodological)*, 50(2), 157–194.
- Liao, H.-J., Liu, J.-G., Wang, L., & Xiang, T. (2019). Differentiable programming tensor networks. *Physical Review X*, 9(3), 031041.
- Liu, Jin Guo, Gao, X., Cain, M., Lukin, M. D., & Wang, S.-T. (2022). *Computing solution space properties of combinatorial optimization problems via generic tensor networks*. *arXiv*. <https://doi.org/10.48550/ARXIV.2205.03718>
- Liu, Jin-Guo, Wang, L., & Zhang, P. (2021). Tropical tensor network for ground states of spin glasses. *Physical Review Letters*, 126(9). <https://doi.org/10.1103/physrevlett.126.090506>
- Markov, I. L., & Shi, Y. (2008). Simulating quantum computation by contracting tensor networks. *SIAM Journal on Computing*, 38(3), 963–981. <https://doi.org/10.1137/050644756>
- Orús, R. (2014). A practical introduction to tensor networks: Matrix product states and projected entangled pair states. *Annals of Physics*, 349, 117–158. <https://doi.org/10.1016/j.aop.2014.06.013>
- Orús, R. (2019). Tensor networks for complex quantum systems. *Nature Reviews Physics*, 1(9), 538–550.

- 120 Pan, F., & Zhang, P. (2021). *Simulating the sycamore quantum supremacy circuits*. <https://arxiv.org/abs/2103.03074>
- 121
- 122 Pearl, J. (1985). Bayesian networks: A model of self-activated memory for evidential reasoning.
- 123 *Proc. Of Cognitive Science Society (CSS-7)*.
- 124 Roa-Villescas, M., Liu, J. G., Wijnings, P. W. A., Stuijk, S., & Corporaal, H. (2023). Scal-
- 125 ing probabilistic inference through message contraction optimization. *2023 Congress in*
- 126 *Computer Science, Computer Engineering, & Applied Computing (CSCE)*.
- 127 Roa-Villescas, M., Wijnings, P. W. A., Stuijk, S., & Corporaal, H. (2022). Partial evaluation in
- 128 junction trees. *2022 25th Euromicro Conference on Digital System Design (DSD)*, 429–437.
- 129 <https://doi.org/10.1109/DSD57027.2022.00064>
- 130 Robeva, E., & Seigal, A. (2019). Duality of graphical models and tensor networks. *Information*
- 131 *and Inference: A Journal of the IMA*, 8(2), 273–288. [https://doi.org/10.1093/imaiai/](https://doi.org/10.1093/imaiai/iaiy009)
- 132 [iaiy009](https://doi.org/10.1093/imaiai/iaiy009)

DRAFT